

Harnessing High-Level Song Descriptors towards Natural Language-Based Music Recommendation

Elena V. Epure, **Gabriel Meseguer Brocal**,
Darius Afchar, Romain Hennequin



DEEZER

NL interfaces for search and recommendations highly in demand


- integrating both search and recommendation
- (ideally) conversational
- promising results with LLMs, but still many challenges




Existing approaches for music

Learn to embed **user input / song description** together with music representation (from **audio** or from **embedded metadata**) using contrastive learning

Multimodal learning



I'd love to hear some more upbeat tropical music with a similar steel drum melody in a higher register.



Dense retrieval



"Island music party" by "Steel drum band"

I'd love to hear some more upbeat tropical music with a similar steel drum melody in a higher register.



"On the nature of daylight" by "Max Richter"

**Still a high semantic
gap to bridge**

**What about using
song descriptors instead?**

Research objective

Given **user input / song description** retrieve / recommend songs with a **matching descriptor set**

I'd love to hear some more
upbeat tropical music with a
similar steel drum melody in
a higher register.

steeldrum, higher register,
amateur recording



Island Cruise Steel Band Presents...

Island Cruise Steel Band
20 tracks | 1 h 23 minutes | 09/08/2011 | 67 fans



Tropical Steel Drum Essentials

Island Steel Drum Players
30 tracks | 1 h 49 minutes | 01/07/2010 | 38 fans

Baselines

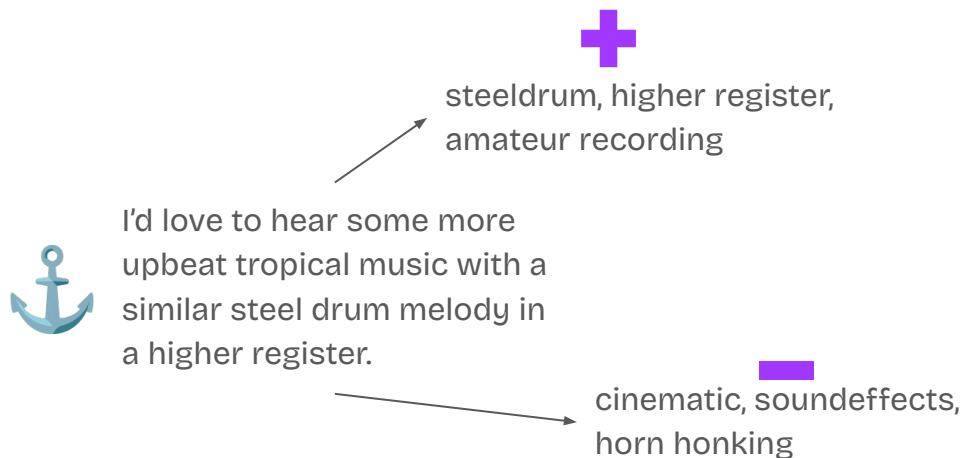
Given **user input / song description** retrieve / recommend songs with a **matching descriptor set**

	sparse vectors	text encoders from multimodal models		PLMs	sentence transformers (dual encoders fine-tuned on sentence similarity)	
	Tf-Idf	CLAP _{text}	TTMR _{text}	BERT	all-MiniLM	msmarco-BERT
MTT	57.7 ± 0.8	13.5 ± 0.3	7.8 ± 0.6	4.8 ± 0.4	33.3 ± 0.6	32.1 ± 0.2
MSD	30.6 ± 2.3	3.4 ± 0.1	5.1 ± 0.1	4.5 ± 0.0	19.5 ± 0.2	20.7 ± 0.1
MC	89.4 ± 0.4	36.5 ± 1.1	19.9 ± 0.2	24.3 ± 0.9	59.9 ± 0.9	66.1 ± 0.2
MC _{reco}	77.7 ± 0.5	27.6 ± 0.4	17.9 ± 1.2	16.3 ± 0.1	48.3 ± 0.4	50.7 ± 1.0

Our proposal

Learn to embed **user input / song descriptions** with **music descriptor sets**

Dense retrieval



Our proposal

Learn to embed **user input / song descriptions** with **music descriptor sets**

Dense retrieval



I'd love to hear some more
upbeat tropical music with a
similar steel drum melody in
a higher register.



steeldrum, higher register,
amateur recording

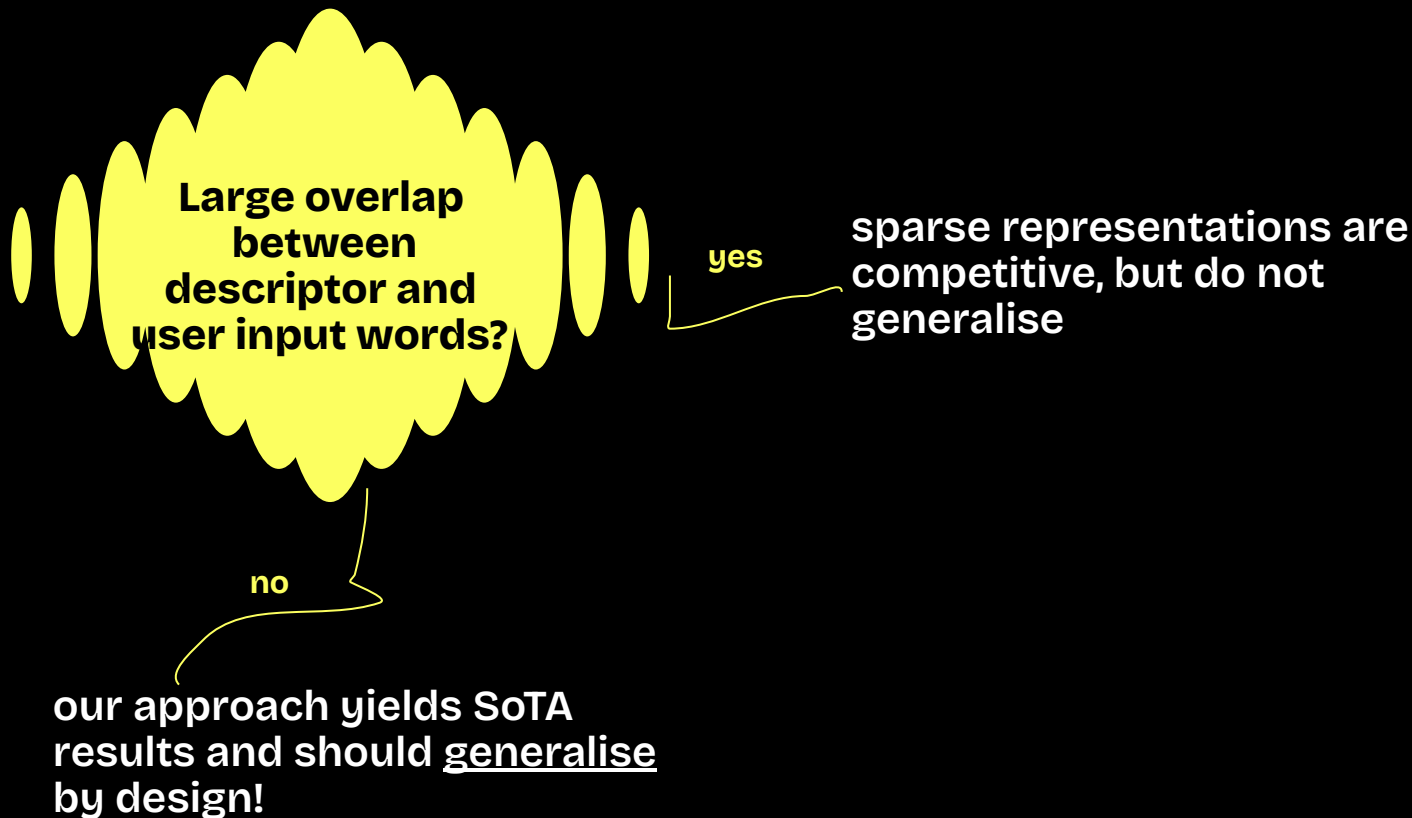


cinematic, soundeffects,
horn honking

Generative Pseudo-labeling Method (GPL)

- Hard negative mining with a music fine-tuned cross-encoder
- Pseudo-labeling of training dataset (anchor, positive, negative)
- Fine-tuning a bi-encoder to mimic the score margin between the positive and negative

Results



Thank you



DEEZER